

# Sel4报告

## sel4组成部分

- CSpace
- 进程间通信 (ipc)
- 虚拟内存 (vspace)
- 中断异常
- 进程控制块

## CSpace

存放能力的空间，内核对象 内存空间均由能力管理

```
#sel4中所有能力均由两个字节组成，存储了具体能力的特定信息。
#sel4将所有的内存全部用untyped_cap管理
block untyped_cap {
    field capFreeIndex 39 //
    padding 18
    field capIsDevice 1
    field capBlockSize 6 //内存大小 2^capBlockSize + capFreeIndex
    field capType 5
    padding 20
    field_high capPtr 39 //指向内存起始地址的指针
},
#sel4具体的内核对象，用于进程间同步通信的端点能力
block endpoint_cap(capEPBadge, capCanGrantReply, capCanGrant, capCanSend,
                  capCanReceive, capEPPtr, capType) {
    field capEPBadge 64
    field capType 5
    field capCanGrantReply 1
    field capCanGrant 1
    field capCanReceive 1
    field capCanSend 1
    padding 16
    field_high capEPPtr 39 //指向内核对象 端点 的指针
}
```

CSpace由cte以链表的形式构成

```
struct cte {
    cap_t cap; //用于存放具体的能力
    mdb_node_t cteMDBNode; //构成一个双向的链表，将同一进程的cte连接起来
};

block mdb_node {
    padding 25
    field_high mdbNext 37
    field mdbRevocable 1
}
```

```

    field mdbFirstBadged 1
    field mdbPrev 64
}

exception_t invokeCNodeRevoke(cte_t *destSlot);
exception_t invokeCNodeDelete(cte_t *destSlot);
exception_t invokeCNodeInsert(cap_t cap, cte_t *srcSlot, cte_t *destSlot);
exception_t invokeCNodeMove(cap_t cap, cte_t *srcSlot, cte_t *destSlot);
exception_t invokeCNodeRotate(cap_t cap1, cap_t cap2, cte_t *slot1,
                             cte_t *slot2, cte_t *slot3);

//srcSlot<=>next
//srcSlot<=>destSlot<=>next
void cteInsert(cap_t newCap, cte_t *srcSlot, cte_t *destSlot)
{
    mdb_node_t srcMDB, newMDB;
    cap_t srcCap;
    bool_t newCapIsRevocable;

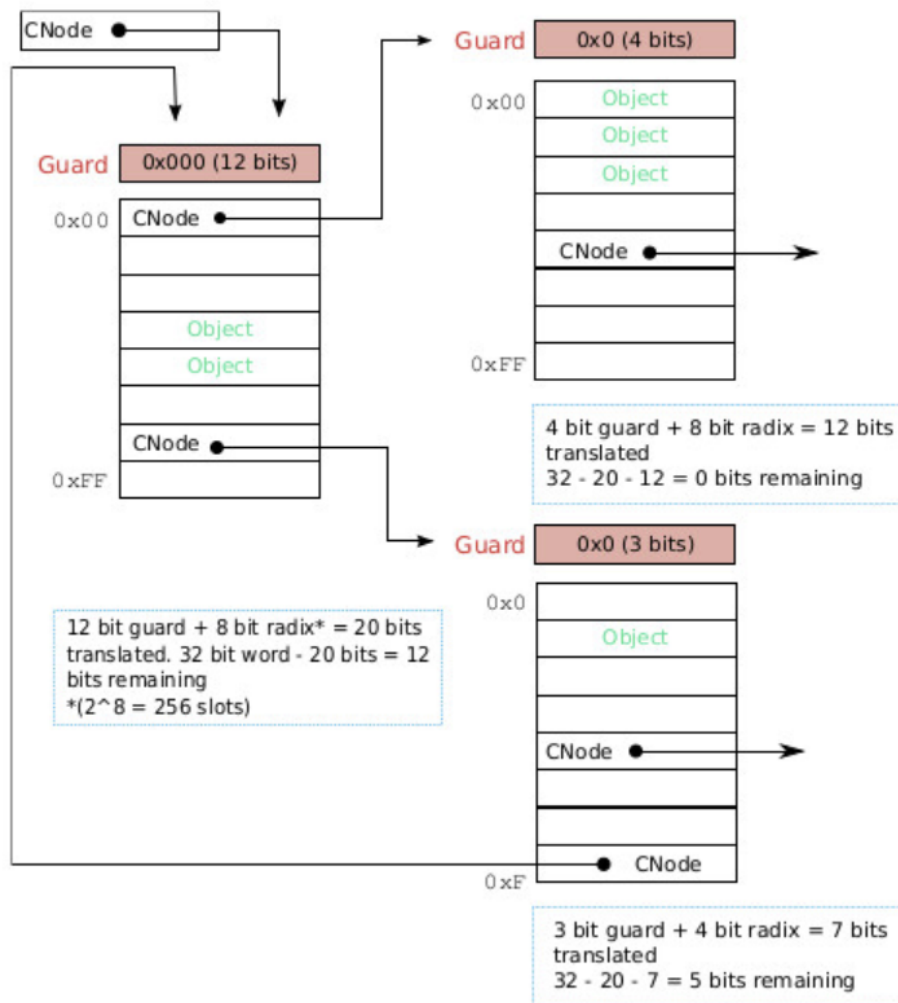
    srcMDB = srcSlot->cteMDBNode;
    srcCap = srcSlot->cap;

    newCapIsRevocable = isCapRevocable(newCap, srcCap);

    //srcSlot<-destSlot->nextSlot
    newMDB = mdb_node_set_mdbPrev(srcMDB, ((word_t)(srcSlot)));
    newMDB = mdb_node_set_mdbRevocable(newMDB, newCapIsRevocable);
    newMDB = mdb_node_set_mdbFirstBadged(newMDB, newCapIsRevocable);
    ;

    destSlot->cap = newCap;
    destSlot->cteMDBNode = newMDB;
    //srcSlot<=>destSlot->nextSlot
    mdb_node_ptr_set_mdbNext(&srcSlot->cteMDBNode, ((word_t)(destSlot)));
    if (mdb_node_get_mdbNext(newMDB)) {
        srcSlot<=>destSlot<=>next
        mdb_node_ptr_set_mdbPrev(
            &((cte_t *) (mdb_node_get_mdbNext(newMDB)))->cteMDBNode,
            ((word_t)(destSlot)));
    }
}

```



```
//cnode查询
block cnode_cap(capCNodeRadix, capCNodeGuardSize, capCNodeGuard,
                capCNodePtr, capType) {
    field capCNodeGuard 64
    field capType 5
    field capCNodeGuardSize 6
    field capCNodeRadix 6
    padding 9
    field_high capCNodePtr 38
}
resolveAddressBits_ret_t resolveAddressBits(cap_t nodeCap, cptr_t capptr, word_t
n_bits)
{
    ...
    while (1) {
        radixBits = cap_cnode_cap_get_capCNodeRadix(nodeCap); //槽号位数
        guardBits = cap_cnode_cap_get_capCNodeGuardSize(nodeCap); //保护位位数
        levelBits = radixBits + guardBits; //当前cnode占用的位数

        capGuard = cap_cnode_cap_get_capCNodeGuard(nodeCap); //当前cnode的保护位

        /* The MASK(wordRadix) here is to avoid the case where
        * n_bits = wordBits (=2^wordRadix) and guardBits = 0, as it
        violates
```

```

        * the C spec to shift right by more than wordBits-1.
        */
        guard = (capptr >> ((n_bits - guardBits) & MASK(wordRadix))) &
MASK(guardBits); //读取要查询的cap的保护位
        if (unlikely(guardBits > n_bits || guard != capGuard)) { //如果两个保护
位不相等就报错
            current_lookup_fault =
                lookup_fault_guard_mismatch_new(capGuard, n_bits,
guardBits);
            ret.status = EXCEPTION_LOOKUP_FAULT;
            return ret;
        }

        if (unlikely(levelBits > n_bits)) { //如果当前cnode占用位数已超过要查询
slot的要求位数则报错
            current_lookup_fault =
                lookup_fault_depth_mismatch_new(levelBits, n_bits);
            ret.status = EXCEPTION_LOOKUP_FAULT;
            return ret;
        }

        offset = (capptr >> (n_bits - levelBits)) & MASK(radixBits);
        slot = CTE_PTR(cap_cnode_cap_get_capCNodePtr(nodeCap)) + offset;

        if (likely(n_bits == levelBits)) { //如果当前cnode占用位数刚好与slot要求查
询的位数想等，则匹配
            ret.status = EXCEPTION_NONE;
            ret.slot = slot;
            ret.bitsRemaining = 0;
            return ret;
        }
        //否则，就是cnode占用位数小于slot查询位数的情况，这说明，当前找到的slot是一个
cnode，仍需向下一层查询

        n_bits -= levelBits;
        nodeCap = slot->cap;

        if (unlikely(cap_get_capType(nodeCap) != cap_cnode_cap)) {
            ret.status = EXCEPTION_NONE;
            ret.slot = slot;
            ret.bitsRemaining = n_bits;
            return ret;
        }
    }
}

```

```

lookupSlot_raw_ret_t lookupSlot(tcb_t *thread, cptr_t capptr)
{
    cap_t threadRoot;
    resolveAddressBits_ret_t res_ret;
    lookupSlot_raw_ret_t ret;

```

```
threadRoot = (((cte_t *) ((word_t) (thread) & ~((1ul << (10)) - 1ul))) +
(tcblCTable)) -> cap;
res_ret = resolveAddressBits(threadRoot, capptr, (1ul << (6)));

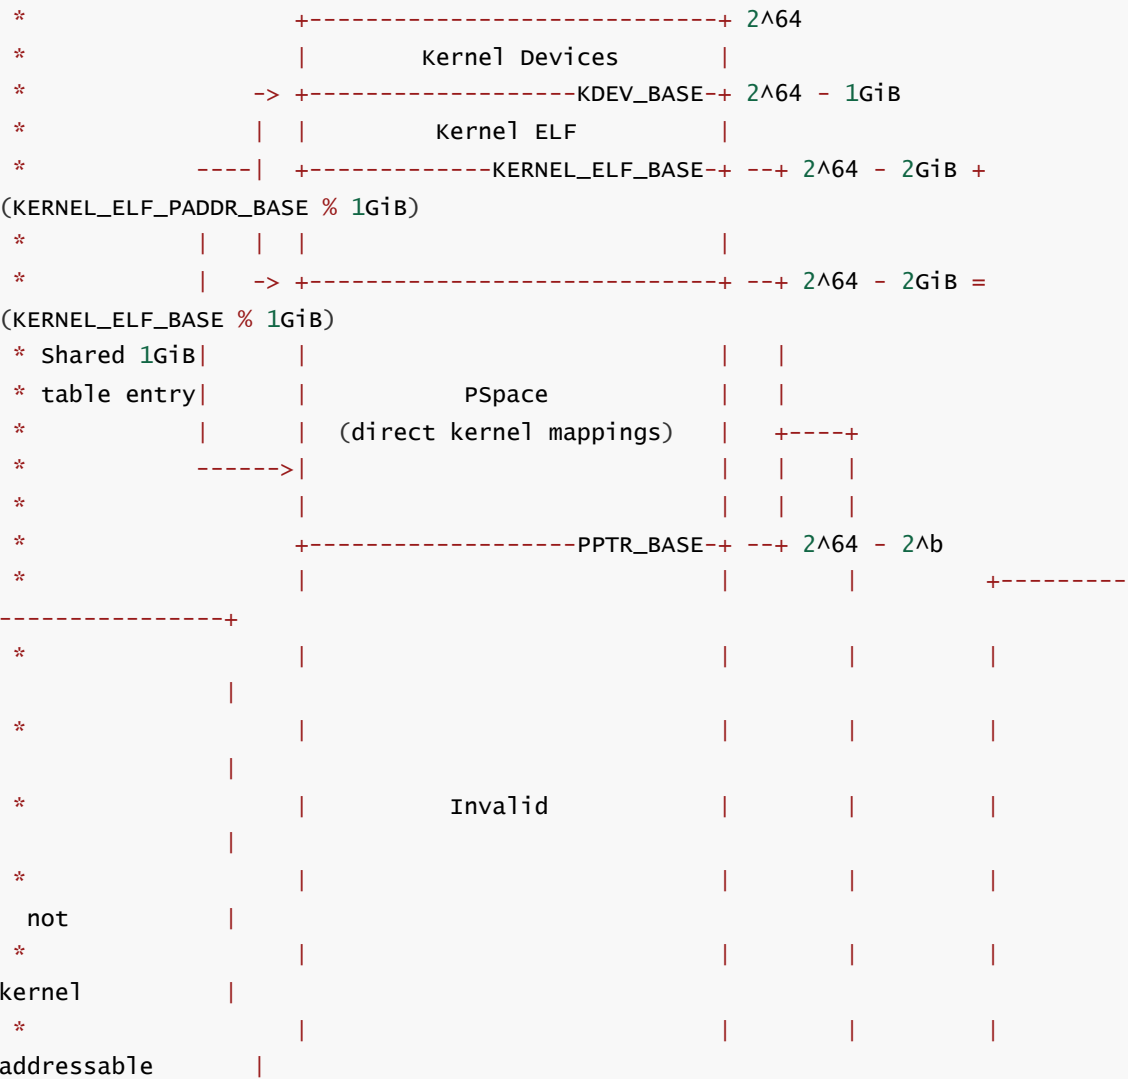

ret.status = res_ret.status;
ret.slot = res_ret.slot;
return ret;
}
```

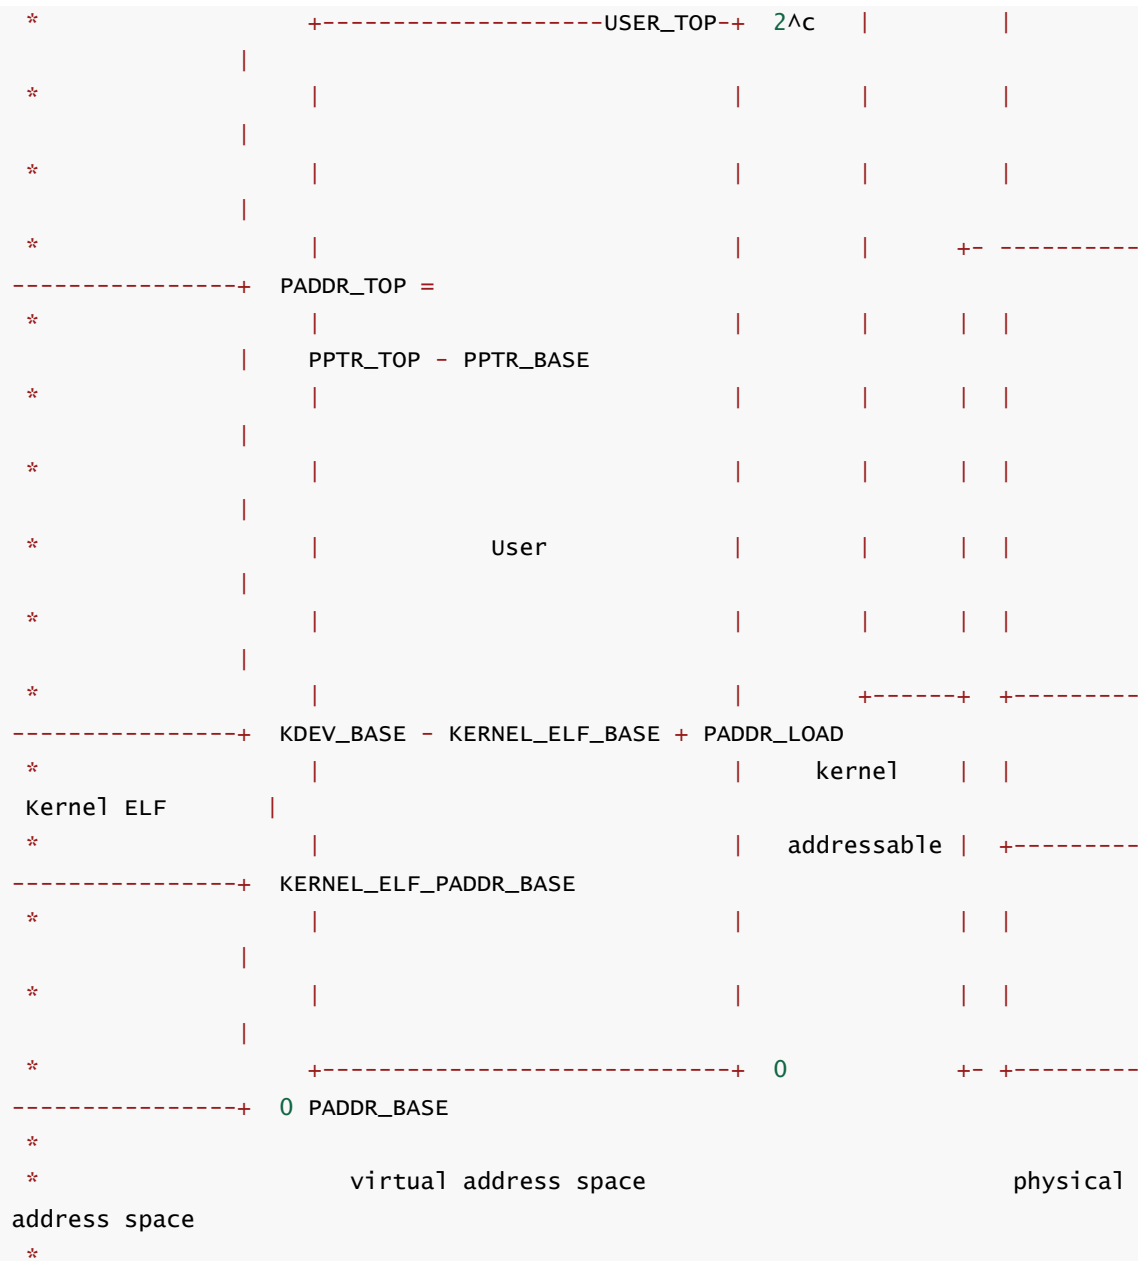
虚拟地址 (VSpace)

sel4的页表与其他操作系统的页表并无太大差别，都采用了sv39的页表设计

Virtual Address (39bits)

地址格式与组成





```

block frame_cap \ page_table_cap {
    field capPTMappedASID 16
    field_high capPTBasePtr 39//页表的基地址
    padding 9
    field capType 5
    padding 19
    field capPTIsMapped 1 //是否被映射
    field_high capPTMappedAddress 39//自身映射的虚拟地址
}

BOOT_CODE void map_it_frame_cap(cap_t vspace_cap, cap_t frame_cap)
{
    pte_t *lvl1pt = PTE_PTR(pptr_of_cap(vspace_cap));
    pte_t *frame_pptr = PTE_PTR(pptr_of_cap(frame_cap));
    vptr_t frame_vptr = cap_frame_cap_get_capFMappedAddress(frame_cap);

```

```

/* We deal with a frame as 4KiB */
lookupPTSlot_ret_t lu_ret = lookupPTSlot(lv1lpt, frame_vptr);
assert(lu_ret.ptBitsLeft == seL4_PageBits);

pte_t *targetSlot = lu_ret.ptSlot;

*targetSlot = pte_new(
    (pptr_to_paddr(frame_pptr) >> seL4_PageBits),
    0, /* sw */
    1, /* dirty (leaf) */
    1, /* accessed (leaf) */
    0, /* global */
    1, /* user (leaf) */
    1, /* execute */
    1, /* write */
    1, /* read */
    1 /* valid */
);
sfence();
}

lookupPTSlot_ret_t lookupPTSlot(pte_t *lv1lpt, vptr_t vptr)
{
    lookupPTSlot_ret_t ret;

    word_t level = CONFIG_PT_LEVELS - 1;
    pte_t *pt = lv1lpt; //当前页表指向1级页表的基址

    ret.ptBitsLeft = PT_INDEX_BITS * level + seL4_PageBits;
    ret.ptSlot = pt + ((vptr >> ret.ptBitsLeft) & MASK(PT_INDEX_BITS)); //获得1级页
    表中对应的槽

    while (isPTEPageTable(ret.ptSlot) && likely(0 < level)) {
        level--;
        ret.ptBitsLeft -= PT_INDEX_BITS;
        pt = getPPtrFromHWPTE(ret.ptSlot); //获得当前对应槽的ppn并将其转换为虚拟地址（下
        一级页表的基址）
        ret.ptSlot = pt + ((vptr >> ret.ptBitsLeft) & MASK(PT_INDEX_BITS)); //获得
        对应的槽
    }

    return ret;
}

BOOT_CODE cap_t create_it_address_space(cap_t root_cnode_cap, v_region_t
it_v_reg)
{
    ...
    lv1lpt_cap =
        cap_page_table_cap_new(
            IT_ASID, /* capPTMappedASID */
            (word_t) rootserver.vspace, /* capPTBasePtr */
            1, /* capPTIsMapped */
            (word_t) rootserver.vspace /* capPTMappedAddress */

```

```

    );

    /* create all n level PT caps necessary to cover userland image in 4KiB pages
    */
    for (int i = 0; i < CONFIG_PT_LEVELS - 1; i++) {

        for (pt_vptr = ROUND_DOWN(it_v_reg.start, RISC_V_GET_LVL_PG_SIZE_BITS(i));
             pt_vptr < it_v_reg.end;
             pt_vptr += RISC_V_GET_LVL_PG_SIZE(i)) {
            if (!provide_cap(root_cnode_cap,
                            create_it_pt_cap(lvl1pt_cap, it_alloc_paging(),
pt_vptr, IT_ASID)))
                ) {
                return cap_null_cap_new();
            }
        }

        ...
    }
    return lvl1pt_cap;
}

```

## IPC

主要通过端点来进行ipc的通信

```

block endpoint {
    field epQueue_head 64
    padding 25
    field_high epQueue_tail 37
    field state 2
}

//tcb结构
struct tcb {
    ...
    word_t tcbIPCBuffer;
    struct tcb *tcbEPNext;
    struct tcb *tcbEPPrev;
    ...
};

block endpoint_cap(capEPBadge, capCanGrantReply, capCanGrant, capCanSend,
                  capCanReceive, capEPPtr, capType) {
    field capEPBadge 64
    field capType 5
    field capCanGrantReply 1
    field capCanGrant 1
    field capCanReceive 1
    field capCanSend 1
    padding 16
    field_high capEPPtr 39 //指向内核对象 端点 的指针
}

```



```

block endpoint_cap(capEPBadge, capCanGrantReply, capCanGrant, capCanSend,
                  capCanReceive, capEPPtr, capType) {
    field capEPBadge 64
    field capType 5
    field capCanGrantReply 1
    field capCanGrant 1
    field capCanReceive 1
    field capCanSend 1
    padding 16
    field_high capEPPtr 39 //指向内核对象 端点 的指针
}

```

//涉及的函数为sendIPC\receiveIPC，均为一个状态机。

```

void sendIPC(bool_t blocking, bool_t do_call, word_t badge,
             bool_t canGrant, bool_t canGrantReply, tcb_t *thread, endpoint_t
             *epptr)

{
    switch (endpoint_ptr_get_state(epptr)) {
    case EPState_Idle:
    case EPState_Send:
        if (blocking) {
            tcb_queue_t queue;

            scheduleTCB(thread); //因为当前线程阻塞在发送端，所以要进行schedule

            queue = ep_ptr_get_queue(epptr); //endpoint对象内部包含一个队列，用于存储当前等待发送的线程
            queue = tcbEPAppend(thread, queue);
            endpoint_ptr_set_state(epptr, EPState_Send);
            ep_ptr_set_queue(epptr, queue);
        }
        break;

    case EPState_Recv: {
        tcb_queue_t queue;
        tcb_t *dest;

        queue = ep_ptr_get_queue(epptr);
        dest = queue.head;

        queue = tcbEPDequeue(dest, queue);
        ep_ptr_set_queue(epptr, queue);

        if (!queue.head) {
            endpoint_ptr_set_state(epptr, EPState_Idle);
        }

        /* Do the transfer */
        doIPCTransfer(thread, ep_ptr, badge, canGrant, dest);
        bool_t replyCanGrant = thread_state_ptr_get_blockingIPCCanGrant(&dest->tcbState);
    }
}

```

```

        setThreadState(dest, ThreadState_Running);
        possibleSwitchTo(dest);

        break;
    }
}

void receiveIPC(tcb_t *thread, cap_t cap, bool_t isBlocking)
{
    endpoint_t *epptr;
    epptr = ((endpoint_t *) (cap_endpoint_cap_get_capEPPtr(cap)));
    switch (endpoint_ptr_get_state(epptr)) {
        case EPState_Idle:
        case EPState_Recv: {
            tcb_queue_t queue;

            if (isBlocking) {
                scheduleTCB(thread);

                /* Place calling thread in endpoint queue */
                queue = ep_ptr_get_queue(epptr);
                queue = tcbEPAppend(thread, queue);
                endpoint_ptr_set_state(epptr, EPState_Recv);
                ep_ptr_set_queue(epptr, queue);
            } else {
                doNBRecvFailedTransfer(thread);
            }
            break;
        }

        case EPState_Send: {
            tcb_queue_t queue;
            tcb_t *sender;
            word_t badge;
            bool_t canGrant;
            bool_t canGrantReply;
            bool_t do_call;

            queue = ep_ptr_get_queue(epptr);
            sender = queue.head;

            queue = tcbEPDequeue(sender, queue);
            ep_ptr_set_queue(epptr, queue);

            if (!queue.head) {
                endpoint_ptr_set_state(epptr, EPState_Idle);
            }

            /* Get sender IPC details */
            badge = thread_state_ptr_get_blockingIPCBadge(&sender->tcbState);
            canGrant =
                thread_state_ptr_get_blockingIPCCanGrant(&sender->tcbState);

```

```

        canGrantReply =
            thread_state_ptr_get_blockingIPCCanGrantReply(&sender-
>tcbState);

        /* Do the transfer */
        doIPCTransfer(sender, ep_ptr, badge,
            canGrant, thread);

        setThreadState(sender, ThreadState_Running);
        possibleSwitchTo(sender);

        break;
    }
}

void doNormalTransfer(tcb_t *sender, word_t *sendBuffer, endpoint_t *endpoint,
    word_t badge, bool_t canGrant, tcb_t *receiver,
    word_t *receiveBuffer)
{
    word_t msgTransferred;
    sel4_MessageInfo_t tag;
    exception_t status;

    tag = messageInfoFromWord(getRegister(sender, msgInfoRegister));

    if (canGrant) { //如果endpoint设置为传递能力，则需将sender的sendbuffer中的能力加入
receiver的Cspace
        status = lookupExtraCaps(sender, sendBuffer, tag);
        if (__builtin_expect(!(status != EXCEPTION_NONE), 0)) {
            current_extra_caps.excaprefs[0] = ((void *)0);
        }
    } else {
        current_extra_caps.excaprefs[0] = ((void *)0);
    }

    msgTransferred = copyMRS(sender, sendBuffer, receiver, receiveBuffer,
        sel4_MessageInfo_get_length(tag));

    tag = transferCaps(tag, endpoint, receiver, receiveBuffer);

    tag = sel4_MessageInfo_set_length(tag, msgTransferred);
    setRegister(receiver, msgInfoRegister, wordFromMessageInfo(tag));
    setRegister(receiver, badgeRegister, badge);
}

```

## Thread

```

//tcb结构
struct tcb {
    /* arch specific tcb state (including context)*/
    arch_tcb_t tcbArch; //保存了全部的寄存器，还多保存了SCAUSE, SSTATUS, SEPC, NextIP（下
一次跳转地址），

```

```

/* Thread state, 3 words */
thread_state_t tcbState;

/* Current fault, 2 words */
sel4_Fault_t tcbFault;//记录当前的错误

/* maximum controlled priority, 1 byte (padded to 1 word) */
prio_t tcbMCP;

/* Priority, 1 byte (padded to 1 word) */
prio_t tcbPriority;

/* Timeslice remaining, 1 word */
word_t tcbTimeslice;

/* Capability pointer to thread fault handler, 1 word */
cptr_t tcbFaultHandler;

/* userland virtual address of thread IPC buffer, 1 word */
word_t tcbIPCBuffer;

/* Previous and next pointers for scheduler queues , 2 words */
struct tcb *tcbSchedNext;
struct tcb *tcbSchedPrev;

/* Previous and next pointers for endpoint and notification queues, 2 words
*/
struct tcb *tcbEPNext;
struct tcb *tcbEPPrev;

};

```

管理不同优先级的tcb的数据结构为位图和队列构成的结构，不同的优先级放入位图的不同位置的队列中。

```

tcb_t* ksSchedulerAction;
tcb_t* SchedulerAction_ResumeCurrentThread=0;
tcb_t* SchedulerAction_ChooseNewThread=1;
void schedule(void)
{
    if (NODE_STATE(ksSchedulerAction) != SchedulerAction_ResumeCurrentThread)
    {
        //如果scheduleraction要求重启当前线程，那就重启当前线程，调度直接完成。
        bool_t was_runnable;
        if (isSchedulable(NODE_STATE(ksCurThread))) {
            //如果当前线程仍然可执行，那就把它加入到ready队列头部
            was_runnable = true;
            SCHED_ENQUEUE_CURRENT_TCB;
        } else {
            was_runnable = false;
        }
    }
}

```

```

        if (NODE_STATE(ksSchedulerAction) == SchedulerAction_ChooseNewThread)
        {
            //如果scheduler要求必须执行新线程，就选择新线程执行。选择的新进程是最高优先级队列中的第一个进程。

            scheduleChooseNewThread();
        } else {
            tcb_t *candidate = NODE_STATE(ksSchedulerAction); //否则，此时
            scheduleraction代表的为竞争线程
            assert(isschedulable(candidate));
            bool_t fastfail =
                NODE_STATE(ksCurThread) == NODE_STATE(ksIdleThread)
                || (candidate->tcbPriority < NODE_STATE(ksCurThread)-
                >tcbPriority);
            if (fastfail &&
                !isHighestPrio(ksCurDomain, candidate->tcbPriority)) { //（竞争线程
                优先级小于当前线程 || 当前线程优先级与初始线程相同）&&竞争线程并非最高优先级。
                SCHED_ENQUEUE(candidate);
                /* we can't, need to reschedule */
                NODE_STATE(ksSchedulerAction) =
                SchedulerAction_ChooseNewThread; //选择新线程而非竞争线程执行。
                scheduleChooseNewThread();
            } else if (was_runnable && candidate->tcbPriority ==
                NODE_STATE(ksCurThread)->tcbPriority) {
                /* we append the candidate at the end of the scheduling queue,
                that way the
                * current thread, that was enqueued at the start of the
                scheduling queue
                * will get picked during chooseNewThread */
                //如果两者优先级相同，仍选择当前线程执行。
                SCHED_APPEND(candidate);
                NODE_STATE(ksSchedulerAction) = SchedulerAction_ChooseNewThread;
                scheduleChooseNewThread();
            } else { //选择新线程执行
                assert(candidate != NODE_STATE(ksCurThread));
                switchToThread(candidate);
            }
        }
    }
    NODE_STATE(ksSchedulerAction) = SchedulerAction_ResumeCurrentThread;
}

```

```

void chooseThread(void)
{
    word_t prio;
    word_t dom;
    tcb_t *thread;

    if (numDomains > 1) {
        dom = ksCurDomain;
    } else {
        dom = 0;
    }

    if (likely(NODE_STATE(ksReadyQueuesL1Bitmap[dom]))) {

```

```

        prio = getHighestPrio(dom);
        thread = NODE_STATE(ksReadyQueues)[ready_queues_index(dom, prio)].head;
        switchToThread(thread);
    } else {
        switchToIdleThread();
    }
}

```

## Sel4 中断处理

Sel4的中断处理并无特色，遵循riscv 基本理念

```

/*traps.S:*/
.global trap_entry
.extern c_handle_syscall
.extern c_handle_fastpath_reply_recv
.extern c_handle_fastpath_call
.extern c_handle_interrupt
.extern c_handle_exception

/*中断异常处理的入口，stvec设置的中断处理地址*/
trap_entry:
csrrw sp, sscratch, sp ;内核栈与用户栈的切换
csrrw t0, sscratch, t0
STORE ra, (0*REGBYTES)(t0)
STORE sp, (1*REGBYTES)(t0) ;在用户占保存返回地址与内核栈的指针
... ;此处省略保存通用寄存器的过程，所有通用寄存器均被保存

csrr x1, sstatus
STORE x1, (32*REGBYTES)(t0)

csrr s0, scause
STORE s0, (31*REGBYTES)(t0) ;读取scause sstatus寄存器

la sp, (kernel_stack_alloc + BIT(CONFIG_KERNEL_STACK_BITS)) :加载内核栈地址

csrr x1, sepc
STORE x1, (33*REGBYTES)(t0);保存sepc

...;接下来根据s0的值判断是否为中断异常，如果是，就调入对应的c函数去处理

```

中断处理部分比较简单，其中 `getActiveIrq` 函数会获得当前被触发的中断，然后进入 `handleInterrupt` 函数具体处理，在此仅展示部分代码：

```

void handleInterrupt(irq_t irq)
{
    switch (intStateIRQTable[(irq)]) {
        ...
        case IRQTimer:

```

```

        timerTick();
        resetTimer();

        break;
    }
    ackInterrupt(irq);
}

```

当完成中断处理后，会进入用户态的函数restore\_user\_context，这一部分用c的内联汇编编写,基本内容与trap.s中内容类似，不做展示。

这一部分的内容的大致流程与中断类似，具体处理函数为c\_handler\_exception

```

void VISIBLE NORETURN c_handle_exception(void)
{
    NODE_LOCK_SYS;

    c_entry_hook();

    word_t scause = read_scause();
    switch (scause) {
    case RISCVInstructionAccessFault:
    case RISCVLoadAccessFault:
    case RISCVStoreAccessFault:
    case RISCVLoadPageFault:
    case RISCVStorePageFault:
    case RISCVInstructionPageFault:
        handleVMFaultEvent(scause);
        break;
    default:
        handleUserLevelFault(scause, 0);
        break;
    }

    restore_user_context();
    UNREACHABLE();
}

```